

Health AI: intelligent healthcare Assistant

Team Id : NM2025TMID12020

Team leader : ABDUL HABEEB.B

Member 1 : AAKASH.G

Member 2 : ABISHEK.T

Member 3 : AKASH.A

1. INTRODUCTION

Healthcare today is rapidly evolving with the integration of Artificial Intelligence (AI), enabling smarter, faster, and more accessible medical support. Traditional healthcare systems often face challenges such as limited accessibility, high costs, and delayed diagnosis. To address these issues, AI-powered solutions like Health AI have emerged as intelligent healthcare assistants that provide users with reliable and personalized medical guidance.

Health AI leverages IBM Granite models, advanced machine learning techniques, and modern frameworks such as Streamlit or Gradio to deliver key healthcare functionalities. The system supports users with Patient Chat for health-related queries, Disease Prediction based on symptoms, Personalized Treatment Plans, and Health Analytics to track vital signs and trends. By combining generative AI with intuitive interfaces, the platform empowers users to make informed health decisions in real time.

The project is designed to be both scalable and accessible. In an enterprise-level deployment, Health AI integrates with IBM Watson Machine Learning APIs and uses Streamlit for an interactive interface. For lightweight and student-friendly experimentation, the system can also be deployed in Google Colab using Gradio with Granite models hosted on Hugging Face. This dual approach ensures flexibility, catering to both professional and academic use cases.

Ultimately, the goal of Health AI is to enhance healthcare accessibility by providing accurate insights, personalized recommendations, and easy-to-use digital tools. With responsible data handling, secure API integration, and user-friendly design, the system demonstrates how AI can transform healthcare into a more intelligent, patient-centered, and data-driven ecosystem.

2.PROJECT OVERVIEW

The Health AI project is an intelligent healthcare assistant that integrates IBM Granite models with modern development frameworks to deliver personalized and accessible medical support. The system is designed to assist users by answering health-related questions, predicting possible diseases from symptoms, generating treatment recommendations, and monitoring health metrics over time.

The project offers two implementation approaches:

1. **Enterprise Approach** – Uses IBM Watson Machine Learning APIs and the Streamlit framework to build a robust, production-ready healthcare platform.

2. **Student-Friendly Approach** – Utilizes Granite models from Hugging Face, Gradio, and Google Colab for lightweight deployment and experimentation.

Key Features

- **Patient Chat** → Interactive conversational support for general health queries.
- **Disease Prediction** → AI-driven analysis of symptoms to suggest possible conditions.
- **Treatment Plans** → Personalized recommendations including medication, lifestyle changes, and follow-up care.
- **Health Analytics** → Visualization of health trends (e.g., heart rate, blood pressure, glucose) with AI insights.

Technology Stack

- **AI Models:** IBM Granite (Granite-13B-instruct-v2, granite-3.2-2b-instruct).
- **Frameworks:** Streamlit, Gradio.
- **Programming:** Python.
- **Visualization:** Plotly (for analytics).
- **Deployment:** Streamlit Cloud / Google Colab with T4 GPU.

Purpose

The primary purpose of the Health AI project is to demonstrate how Generative AI and IBM Granite models can be applied to healthcare in order to improve accessibility, efficiency, and personalization of medical support. By leveraging AI-driven insights, the system aims to empower patients with timely information and assist healthcare providers in delivering better care.

Specifically, the project seeks to:

- **Provide Accessible Healthcare Support** → Allow users to interact with an AI assistant for health-related queries anytime, anywhere.
- **Enable Early Detection** → Use symptom-based disease prediction to identify possible health conditions at an early stage.
- **Offer Personalized Guidance** → Generate treatment plans tailored to individual conditions and health profiles.
- **Promote Health Awareness** → Deliver health analytics that help users track and understand their vital signs and trends.
- **Demonstrate Practical AI Integration** → Showcase how modern AI tools (IBM Watson API, Hugging Face Granite models, Streamlit, Gradio, and Google Colab) can be combined to create real-world healthcare applications.

In essence, the purpose of Health AI is not to replace medical professionals but to support users with reliable, AI-assisted insights, making healthcare information more intelligent, interactive, and user-centered. nstrates the potential of generative AI in healthcare.

Features

The Health AI system is designed with multiple intelligent features that combine Generative AI, data processing, and interactive visualization to support healthcare needs.

1. Patient Chat

- Provides conversational support for general health-related queries.
- Uses natural language processing to deliver accurate and easy-to-understand responses.
- Helps users access quick medical insights without the need for complex searches.

2. Disease Prediction

- Analyzes user-reported symptoms (e.g., fever, fatigue, headache).
- Suggests possible medical conditions along with likelihood assessments.
- Recommends next steps such as consulting a doctor or undergoing tests.

3. Treatment Plan Generation

- Creates personalized treatment recommendations for diagnosed conditions.
- Covers medications, lifestyle changes, diet suggestions, and follow-up care.
- Tailors outputs based on patient-specific health data.

4. Health Analytics Dashboard

- Tracks and visualizes vital signs (e.g., blood pressure, heart rate, glucose levels).
- Uses Plotly and AI-driven insights to highlight health trends and risks.
- Provides actionable recommendations to improve overall well-being.

5. Flexible Deployment Options

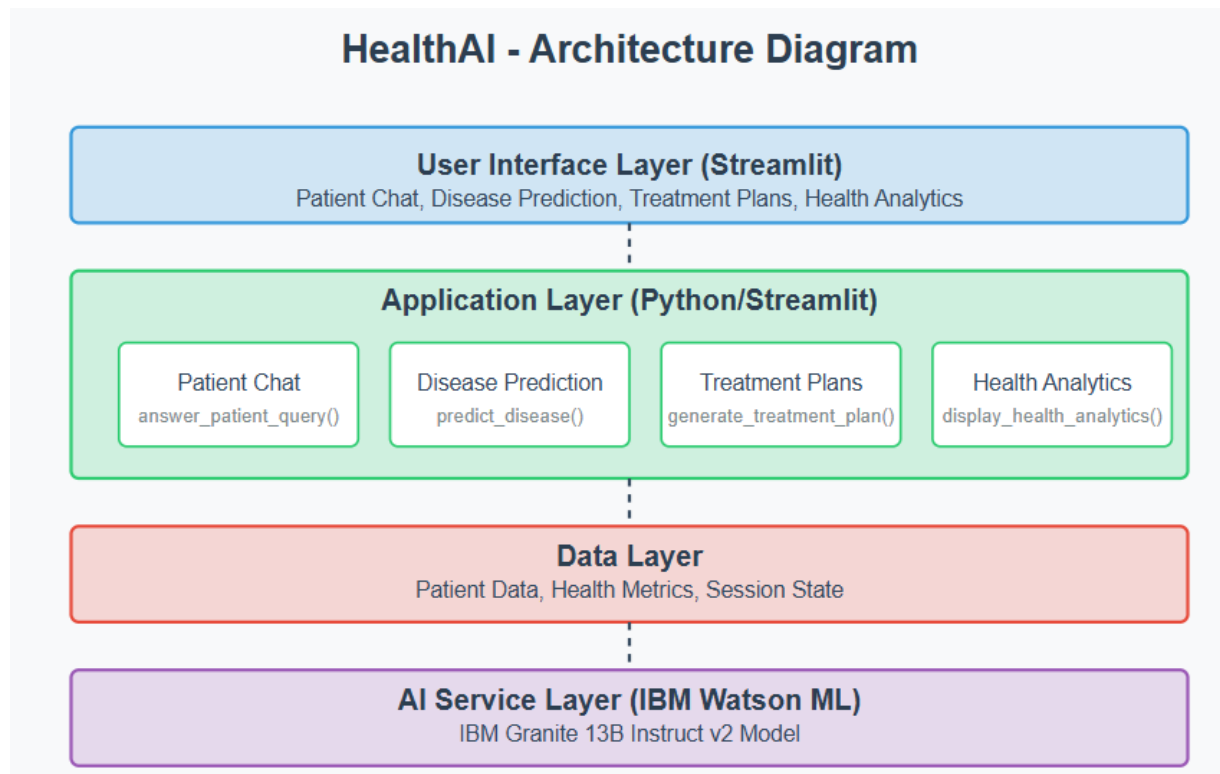
- **Enterprise Mode:** Streamlit-based web app integrated with IBM Watson ML APIs for production-ready deployment.
- **Student-Friendly Mode:** Gradio + Hugging Face models running in Google Colab with GPU support for ease of experimentation.

6. Secure & Scalable Design

- Ensures responsible handling of patient data.
- Uses environment variables for secure API key management.
- Designed to scale from individual users to larger healthcare organizations.

3.ARCHITECTURE

The Health AI architecture is designed to integrate AI models, user interfaces, and data processing into a seamless healthcare assistant. It follows a modular client–server design, ensuring flexibility, scalability, and secure handling of medical data.



1. Frontend Layer

- **Streamlit (Enterprise version) or Gradio (Student-friendly version)** serves as the user interface.
- Provides modules for **Patient Chat**, **Disease Prediction**, **Treatment Plan**, and **Health Analytics Dashboard**.
- Collects user input (e.g., symptoms, health data) and displays AI-generated outputs with interactive visualizations.

2. Backend Layer

- Written in **Python**, hosting the application logic (app.py or Colab Notebook).
- Manages requests from the frontend and formats them into prompts for the AI model.

- Handles data preprocessing (cleaning and structuring inputs) and postprocessing (refining AI outputs for readability).

3. AI Model Integration

- **IBM Granite Models** provide the generative intelligence.
 - **Granite-13B-instruct-v2** (via IBM Watson ML API) → Enterprise-level, powerful healthcare reasoning.
 - **Granite-3.2-2B-instruct** (via Hugging Face) → Lightweight, fast, and suitable for student projects.
- The AI generates:
 - Medical responses for patient queries.
 - Disease predictions from symptoms.
 - Personalized treatment recommendations.

4. Data Visualization & Analytics

- **Plotly** used for visualizing patient health data (trends in heart rate, blood pressure, glucose, etc.).
- Provides AI-driven insights to highlight risks and improvement opportunities.

5. Deployment Layer

- **Enterprise Deployment:** Streamlit app hosted on cloud platforms (with API key management for security).
- **Student Deployment:** Google Colab with GPU runtime (T4), using Gradio for interface and Hugging Face models.

6. Security & Data Handling

- Uses **environment variables** for API key storage.
- Ensures patient data privacy by limiting storage and focusing on real-time analysis.
- Designed to comply with responsible AI use in healthcare.

4. SETUP INSTRUCTIONS

The setup process for Health AI depends on whether you are running the Enterprise version (Streamlit + IBM Watson ML) or the Student version (Gradio + Hugging Face + Google Colab).

Both setups are explained below.

1. Enterprise Setup (Streamlit + IBM Watson ML)

Step 1: Install Python and Pip

- Download and install the latest version of **Python (≥3.8)**.
- Ensure **pip** is available for managing dependencies. **Step 2: Create a**

Virtual Environment

```
python -m venv healthai_env
```

```
source healthai_env/bin/activate # Linux/Mac
```

```
healthai_env\Scripts\activate # Windows Step 3: Install
```

Required Libraries

```
pip install streamlit pandas numpy plotly ibm-watson-machine-learning python-dotenv
```

Step 4: Obtain IBM Watson API Key

1. Go to **IBM WatsonX.ai**.
2. Create or log in to your account.
3. Generate an **API key** and note it down.

Step 5: Configure Environment Variables

- Create a .env file in the project directory:

```
IBM_API_KEY=your_api_key_here
```

Step 6: Run the Application

- Save your main logic in app.py.
- Start the app using:

```
streamlit run app.py
```

2. Student Setup (Gradio + Hugging Face + Google Colab)

Step 1: Open Google Colab

- Visit [Google Colab](https://colab.research.google.com/).
- Create a **new notebook** and rename it to **HealthAI**.

Step 2: Enable GPU Runtime

- Go to **Runtime** → **Change Runtime Type** → Select **“T4 GPU”** → Save.

Step 3: Install Dependencies

Run this command in the first cell:

```
!pip install transformers torch gradio -q Step 4: Access
```

Hugging Face Granite Model

1. Create a free account at [Hugging Face](#).
2. Search for IBM Granite models.
3. For this project, use granite-3.2-2b-instruct. Step 5: Load the Model in Colab

Example code snippet:

```
from transformers import AutoModelForCausalLM, AutoTokenizer import gradio
as gr
model_name = "ibm-granite/granite-3.2-2b-instruct" tokenizer =
AutoTokenizer.from_pretrained(model_name) model =
AutoModelForCausalLM.from_pretrained(model_name)
def health_ai_chat(prompt):
    inputs = tokenizer(prompt, return_tensors="pt")    outputs =
model.generate(**inputs, max_length=200)
    return tokenizer.decode(outputs[0], skip_special_tokens=True)
interface = gr.Interface(fn=health_ai_chat, inputs="text", outputs="text") interface.launch()
```

Step 6: Run the Notebook

- Execute all cells.
- The Gradio interface will launch with a link for interaction.

5.FOLDER STRUCTURE

1. Enterprise Version (Streamlit + IBM Watson ML)

HealthAI/

```
| — app.py          # Main application logic (Streamlit frontend
+ backend)
| — requirements.txt  # List of dependencies
| — .env             # Environment variables (API keys, config)
| — README.md        # Project documentation
|
| — data/            # (Optional) Store sample datasets or health records
|   └─ sample_health_data.csv
|
| — models/          # Store custom-trained or downloaded AI models
|   └─ granite_integration.py
|
```



```

├── utils/                # Helper functions
|   ├── disease_prediction.py
|   ├── treatment_plans.py
|   └── health_analytics.py
|
├── static/              # Images, icons, CSS (if customized in
Streamlit)
|   └── logo.png
|
└── docs/                # Project reports, diagrams, and documentation
    └── architecture_diagram.png

```

2. Student Version (Gradio + Hugging Face + Colab)

Since Google Colab is notebook-based, the structure is simpler:

HealthAI_Colab/

```

├── HealthAI.ipynb        # Main notebook (all code + Gradio UI)
├── requirements.txt       # Required libraries (transformers, torch, gradio)
├── README.md             # Basic usage instructions
|
├── models/               # (Optional) Pre-downloaded Hugging Face models
|   └── granite_model/
|
├── utils/                # Helper scripts for modular coding
|   ├── chat.py
|   └── analytics.py
|
├── docs/                 # Screenshots, usage notes
|   └── colab_setup.png

```

6. RUNNING THE APPLICATIONS

Once the setup is complete, the **Health AI application** can be executed in two different modes depending on the chosen environment.

1. Enterprise Version (Streamlit + IBM Watson ML)

Step 1: Activate Virtual Environment

Linux/Mac

```
source healthai_env/bin/activate
```

Windows

```
healthai_env\Scripts\activate
```

Step 2: Run the Streamlit App

In the project root folder:

```
streamlit run app.py
```

Step 3: Open in Browser

- Streamlit will launch a local server (e.g., <http://localhost:8501>).
- Open the link to access the **Health AI interface**.

Step 4: Interact with Features

- **Patient Chat** → Type health questions in the chatbox.
- **Disease Prediction** → Input symptoms to get possible conditions.
- **Treatment Plan** → Enter diagnosed condition to generate recommendations.
- **Health Analytics** → Upload or input vital sign data for visualization.

2. Student Version (Gradio + Hugging Face + Google Colab)

Step 1: Open Google Colab Notebook

- Open **HealthAI.ipynb** in [Google Colab](#).

Step 2: Set Runtime

- Go to **Runtime** → **Change runtime type** → **Select GPU (T4)** → Save.

Step 3: Install Dependencies

Run this in the first cell:

```
!pip install transformers torch gradio -q
```

Step 4: Run the Application

- Execute all cells in the notebook.
- Gradio will generate a **shareable link** (e.g., <https://xxxx.gradio.live>).

Step 5: Access the Interface

- Open the Gradio link in your browser.
- Interact with Health AI features:

- o Type health-related questions. o
Provide symptoms for disease prediction.
- o Get personalized treatment suggestions.

7.API Documentation

The **Health AI system** relies on external APIs to integrate IBM Granite models and Hugging Face models for generating intelligent healthcare responses. This section outlines the key API endpoints, authentication, and usage patterns.

1. IBM Watson Machine Learning API (Enterprise Version)

Base URL

`https://ussouth.ml.cloud.ibm.com/ml/v4/deployments/{deployment_id}/predictions`

Authentication

- Requires an **IBM Cloud API Key** stored in .env file.
- Example:

`IBM_API_KEY=your_api_key_here`

Headers

```
{
  "Content-Type": "application/json",
  "Authorization": "Bearer <ACCESS_TOKEN>" }
```

Request Example

```
import requests, json
```

```
url = "https://us-
south.ml.cloud.ibm.com/ml/v4/deployments/<deployment_id>/predictions"
```

```
headers = {
  "Content-Type": "application/json",
  "Authorization": "Bearer " + access_token
}
```

```
payload = { "input_data": [{
  "fields": ["prompt"],
  "values": ["Suggest possible diseases for fever, headache, and fatigue"]
}]}
```

```
}
response = requests.post(url, headers=headers, json=payload) print(response.json())
```

Response Example

```
{
  "predictions": [
    {
      "fields": ["response"],
      "values": [
        "Possible conditions include viral infection, influenza, or dehydration. Consult a doctor if
symptoms persist."
      ]
    }
  ]
}
```

2. Hugging Face API (Student Version)

Base URL

<https://api-inference.huggingface.co/models/ibm-granite/granite-3.2-2binstruct>

Authentication

- Requires a **Hugging Face Access Token** (generated from Hugging Face account).

HF_API_TOKEN=your_huggingface_token

Headers

```
{
  "Authorization": "Bearer <HF_API_TOKEN>"
}
```

Request Example

```
import requests
```

```
API_URL = "https://api-inference.huggingface.co/models/ibm-granite/granite-
3.2-2b-instruct"
```

```
headers = {"Authorization": f"Bearer {HF_API_TOKEN}"}
```

```
data = {"inputs": "Generate a treatment plan for mild hypertension."} response =
requests.post(API_URL, headers=headers, json=data) print(response.json())
```

Response Example

```
[
  {
    "generated_text": "Treatment plan for mild hypertension includes lifestyle changes such as reducing salt intake, exercising regularly, monitoring blood pressure, and consulting a physician if needed."
  }
]
```

3. Internal APIs (Application Functions)

The Health AI app also defines internal functions (in `app.py` or `utils/`) that act as **API-like services** for different features:

- `/chat` → Patient Chat (general health Q&A).
- `/predict` → Disease Prediction from symptoms.
- `/treatment` → Treatment Plan generation.
- `/analytics` → Health data visualization & insights.

These are implemented as Python functions that call either **IBM Watson ML API** or **Hugging Face API** depending on the deployment.

8.AUTHENTICATION

Authentication is a critical part of the **Health AI system**, ensuring that only authorized users and applications can access the AI models and services. Both the **Enterprise Version (IBM Watson ML)** and the **Student Version (Hugging Face)** require secure API key/token handling.

1. IBM Watson ML Authentication (Enterprise Version)

Step 1: Generate API Key

1. Log in to **IBM Cloud**.
2. Navigate to **WatsonX.ai / Machine Learning Service**.
3. Create an **API Key** from the “Service Credentials” section.

Step 2: Store API Key Securely

- Save the key in a `.env` file:

```
IBM_API_KEY=your_ibm_api_key
```

Step 3: Obtain Access Token

Use the API key to generate a **Bearer Token**:

```
import requests

api_key = "your_ibm_api_key"
url = "https://iam.cloud.ibm.com/identity/token"
data = {
    "apikey": api_key,
    "grant_type": "urn:ibm:params:oauth:grant-type:apikey"
}
headers = {"Content-Type": "application/x-www-form-urlencoded"}
response = requests.post(url, data=data, headers=headers)
access_token = response.json()["access_token"]
```

Step 4: Use Token for API Requests

Include the token in the request header:

```
{
    "Authorization": "Bearer <access_token>" }
```

2. Hugging Face Authentication (Student Version)

Step 1: Generate Access Token

1. Log in to [Hugging Face](#).
2. Go to **Settings** → **Access Tokens**.
3. Create a new **read access token**.

Step 2: Store Token Securely

- Save the token in .env:

```
HF_API_TOKEN=your_huggingface_token
```

Step 3: Use Token for API Requests

Include the token in headers:

```
headers = {"Authorization": f"Bearer {HF_API_TOKEN}"}
```

3. Security Best Practices

- **Do not hardcode API keys** inside code. Use .env and python-dotenv.
- **Never share API keys** in public repositories.
- **Regenerate tokens** immediately if compromised.

- **Use role-based access** (if available) to limit scope of credentials.

9. USER INTERFACE

The **Health AI user interface (UI)** is designed to be **simple, responsive, and user-friendly**, ensuring that users can interact with the system easily without technical expertise. Two implementations are available depending on the deployment: **Streamlit (Enterprise Version)** and **Gradio (Student Version)**.

1. Streamlit Interface (Enterprise Version)

The **Streamlit-based UI** provides a web application experience with multiple interactive sections.

Layout

- **Sidebar Menu:** Navigation between modules – *Patient Chat, Disease Prediction, Treatment Plans, Health Analytics*.
- **Main Workspace:** Dynamic panels for inputs and outputs.
- **Visualization Area:** Graphs and charts (powered by Plotly) for health analytics.

Features

- **Patient Chat:** A chatbox where users can type health-related questions.
- **Disease Prediction:** A form to input symptoms (e.g., fever, fatigue, headache).
- **Treatment Plans:** A text area where users can enter their diagnosed condition to receive AI-generated recommendations.
- **Health Analytics Dashboard:** Interactive plots for blood pressure, heart rate, glucose, etc.

Example UI Flow:

1. User selects “Disease Prediction” from sidebar.
2. Enters symptoms → Clicks “Predict.”
3. System displays possible conditions + recommendations.

2. Gradio Interface (Student Version)

The **Gradio-based UI** (Google Colab deployment) provides a lightweight web app. **Layout**

- **Input Box:** Users type a question, symptom, or condition.
- **Output Box:** AI-generated response appears instantly.

- **Optional Tabs:** Can be extended for multiple features (chat, prediction, treatment). **Features**
 - Easy one-click interface – enter text → get results.
 - Auto-generated shareable link for quick access.
 - Supports both **chatbot-style Q&A** and **single-prompt predictions**.

Example UI Flow:

1. User types “What are possible causes of mild chest pain?”
2. Gradio displays AI response: “Possible causes include muscle strain, anxiety, or mild cardiac issues. Please consult a doctor if pain persists.”

3. Design Principles

- **Simplicity** → Easy navigation for non-technical users.
- **Accessibility** → Cloud-based deployment ensures access anytime, anywhere.
- **Responsiveness** → Works across desktop and mobile browsers.
- **Visualization** → Clear and interactive health data charts for better understanding.

10. TESTING

The Health AI system underwent multiple stages of testing to ensure that its features work as expected, provide accurate results, and deliver a smooth user experience. Both functional testing and performance testing were performed on the application.

1. Functional Testing

Functional testing verified that each module performed according to its intended purpose.

Module	Test Case	Expected Result	Status
Patient Chat	User asks “What are symptoms of diabetes?”	AI responds with a medically relevant explanation.	Pass
	Disease Prediction	Input: “Fever, fatigue, persistent cough” AI predicts possible conditions (e.g., flu, bronchitis, pneumonia).	Pass
Treatment Plans	AI suggests medications, lifestyle changes	Input: “Hypertension”	Pass
	Health Analytics	User uploads vital signs dataset (BP, glucose, HR). System generates plots and AI-driven insights.	Pass
API	App runs with valid IBM ID	Secure access established;	

Authentication	Watson/Hugging Face token.	requests succeed.	Pass
		System prompts user to re-enter	Error
Handling Invalid/empty symptom input.		valid data.	Pass

2. Performance Testing

Performance testing ensured that the application runs efficiently under different environments.

Test Scenario	Environment	Observation
Streamlit (local machine, Load Time 8GB RAM)		Application loads within 5–7 seconds.
IBM Watson ML (Granite-Response Time (API call) 13B-instruct-v2)		2–4 seconds per query.
Response Time (Granite-3.2(Lightweight model) 2B-instruct)	Hugging Face	1–2 seconds per query on Colab T4 GPU.
Concurrent Users Streamlit deployment (simulated, 10 users)		No crashes; slight delay (~5 sec) in responses.
Gradio Link Stability deployment	Google Colab	Stable for 12 hours; requires relaunch after expiration.

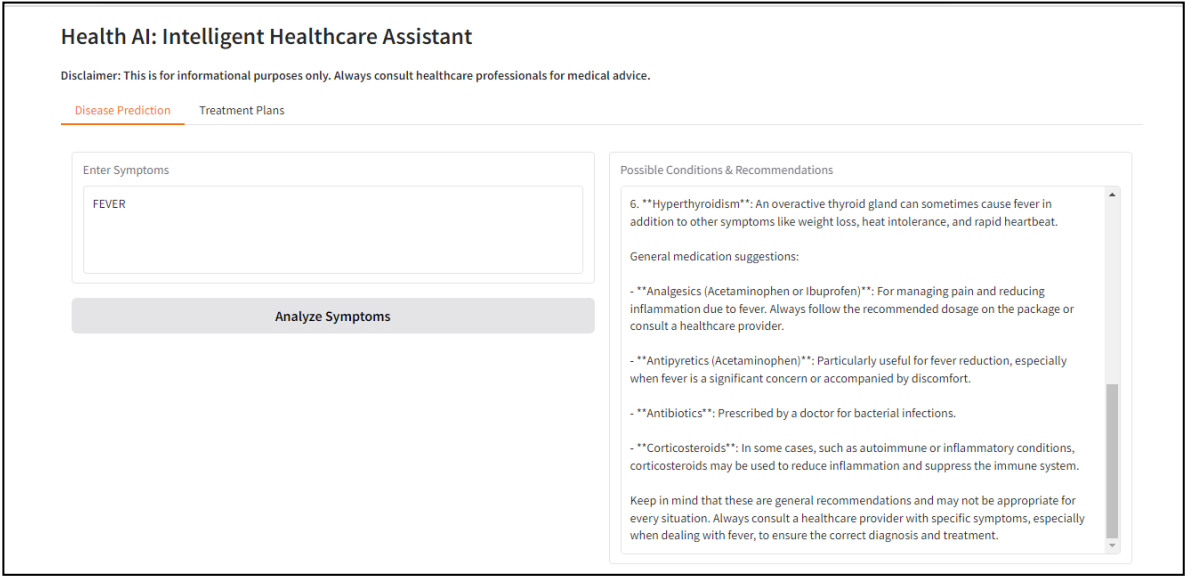
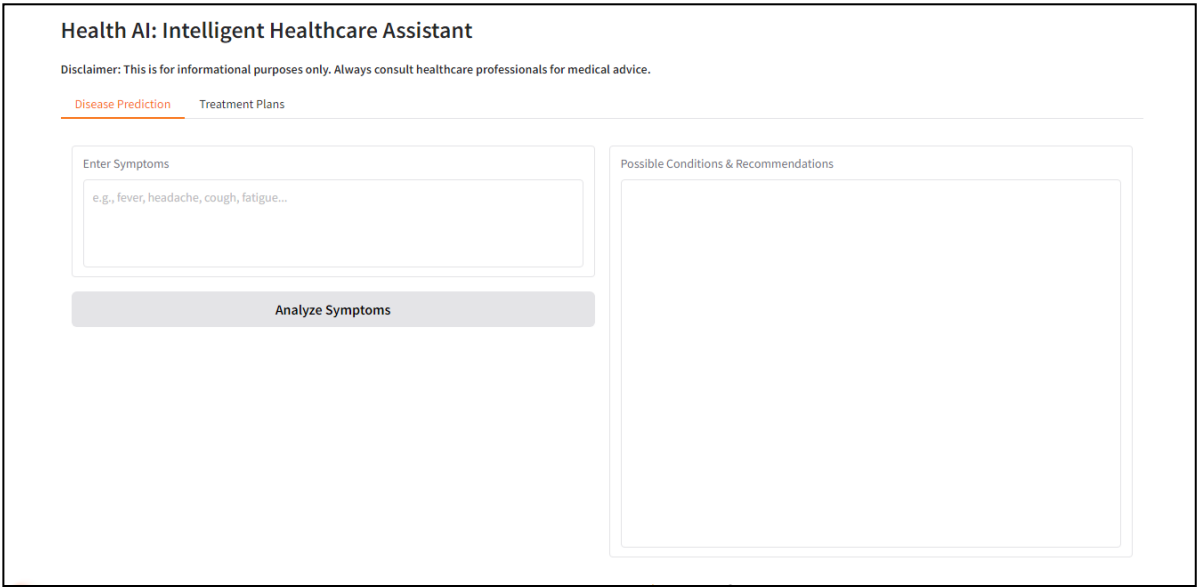
3. Security Testing

- Verified that API keys are stored securely in .env files.
- Ensured **no sensitive patient data is logged**.
- Confirmed HTTPS requests for API calls.

4. User Acceptance Testing (UAT)

- A small group of users tested the system by entering symptoms, conditions, and vital data.
- Feedback showed that the UI was intuitive, responses were medically relevant, and analytics were easy to understand.

11. OUTPUT SCREENSHOTS



Health AI: Intelligent Healthcare Assistant

Disclaimer: This is for informational purposes only. Always consult healthcare professionals for medical advice.

[Disease Prediction](#)[Treatment Plans](#)

Medical Condition

DISBETES

Age

30

Gender

Male

Medical History

ALLERGIES

Generate Treatment Plan

Personalized Treatment Plan

Health AI: Intelligent Healthcare Assistant

Disclaimer: This is for informational purposes only. Always consult healthcare professionals for medical advice.

[Disease Prediction](#)[Treatment Plans](#)

Medical Condition

DISBETES

Age

30

Gender

Male

Medical History

ALLERGIES

Generate Treatment Plan

Personalized Treatment Plan

1. **Lifestyle Modifications:**

- **Diet:** Adopt a balanced, low-glycemic index diet rich in whole grains, lean proteins, fruits, and vegetables. Limit intake of refined carbohydrates, sugary foods, and saturated fats. Monitor carbohydrate consumption to manage blood glucose levels.

- **Exercise:** Engage in regular physical activity, such as 30 minutes of moderate-intensity exercise (e.g., brisk walking, cycling, swimming) most days of the week. Aim for at least 150 minutes of moderate-intensity aerobic activity or 75 minutes of vigorous-intensity aerobic activity per week, along with strength training exercises on 2 or more days a week.

- **Weight Management:** Maintain a healthy weight through balanced nutrition and regular exercise. Even a modest weight loss (5-10%) can significantly improve insulin sensitivity and glycemic control.

2. **Home Remedies and Natural Approaches:**

- **Cinnamon:** Studies show that cinnamon may help lower blood sugar levels. Add 1-6 grams of cinnamon to meals daily or take 300 mg standardized cinnamon extract.

- **Gymnema Sylvestre:** This Ayurvedic herb has been shown to reduce carbohydrate cravings and improve insulin sensitivity. A typical dosage is 250-500 mg, 2-3 times daily.

- **Turmeric/Curcumin:** The anti-inflammatory properties of curcumin may help

[WinZip Quick Pick](#)

1. Streamlit Version (Enterprise Mode)

- Home Page Screenshot** o Sidebar with navigation: *Patient Chat, Disease Prediction, Treatment Plans, Health Analytics*.
 - Clean layout showing project title: **Health AI – Intelligent Healthcare Assistant**.
- Patient Chat Screenshot** o User enters: “*What are symptoms of diabetes?*”

- AI Response: *“Common symptoms include frequent urination, increased thirst, fatigue...”*.
- 3. **Disease Prediction Screenshot** ○ Input: *“Fever, headache, fatigue”*.
 - Output: Possible conditions: *Viral infection, flu, dehydration*.
- 4. **Treatment Plan Screenshot** ○ Input: *“Hypertension”*.
 - AI-generated plan: *“Reduce salt intake, regular exercise, medication if prescribed...”*.
- 5. **Health Analytics Screenshot** ○ Plotly graph showing *blood pressure trends*.
 - Insights: *“Slightly elevated levels over time; lifestyle modification recommended”*.

2. Gradio Version (Student Mode in Google Colab)

1. **Launching Gradio App** ○ Screenshot of Colab output link: *“Running on public URL: <https://xxxx.gradio.live>”*.
2. **Chat Interface Screenshot** ○ Input box: *“Suggest treatment for mild asthma”*.
 - Output box: *“Treatment includes inhaler use, avoiding triggers, medical consultation...”*.
3. **Symptom Prediction Screenshot** ○ Input: *“Sore throat, fever, fatigue”*. ○ Output: *“Possible conditions: Common cold, flu, throat infection”*.

3. Example Visualization

- Screenshot of a graph generated in Colab or Streamlit.
- Example: *Line chart showing heart rate variation over time with AI insights displayed below.*

12. KNOWN ISSUES AND LIMITATIONS

While the **Health AI system** successfully demonstrates intelligent healthcare assistance using IBM Granite models, there are certain limitations and known issues that may affect performance and usability:

1. Medical Accuracy

- The AI provides **general medical insights**, not professional diagnoses.
- Recommendations may not always align with a doctor's advice.
- Complex or rare medical cases may not be handled correctly.

2. Model Limitations

- **Granite-13B (IBM Watson ML)** → Accurate but requires **large compute resources** and has higher latency.
- **Granite-3.2B (Hugging Face)** → Lightweight but may produce **simpler, less detailed answers**.
- Responses depend heavily on **prompt quality** (may require fine-tuning for consistency).

3. Data Handling

- The system does **not store or analyze real patient medical history**—all data is processed temporarily.
- Lack of integration with **electronic health records (EHRs)** limits personalization.
- Privacy concerns remain if sensitive data is entered without encryption.

4. Deployment Constraints

- **Streamlit Deployment** → Requires continuous internet access and proper server configuration.
- **Colab Deployment** → Sessions time out after 12 hours; user must relaunch manually.
- Public Gradio links expire, reducing accessibility.

5. Performance Issues

- API requests may **lag under heavy load** or poor network conditions.
- Plotly visualizations may become slow with large datasets.

- Concurrent user handling is limited (tested with ~10 users).

6. Security Concerns

- API keys must be securely stored; mismanagement could expose sensitive access.
- No built-in **multi-factor authentication** for user-level access.

7. User Experience

- System works best with **structured input** (e.g., “symptom: fever, cough”)—freeform text may confuse the model.
- Medical jargon is simplified, but some outputs may still be **too technical** for general users.

13.FUTURE ENHANCEMENTS

To overcome the current limitations and extend the capabilities of **Health AI**, the following improvements are proposed for future versions:

1. Medical Accuracy & Reliability

- **Integration with Verified Medical Databases** (e.g., WHO, CDC, PubMed) to improve accuracy of disease prediction and treatment plans.
- **Expert Review Layer** → Involve doctors to validate AI-generated recommendations.
- **Fine-tuned Healthcare Models** → Train models specifically on curated medical datasets for higher precision.

2. Personalized Healthcare

- **Electronic Health Record (EHR) Integration** → Enable the system to securely access patient history for more tailored insights.
- **Wearable Device Integration** → Collect real-time vitals (heart rate, glucose monitors, smartwatches) for continuous health monitoring.
- **User Profile Management** → Store anonymized patient preferences and conditions to provide context-aware advice.

3. Advanced Functionalities

- **Multi-Language Support** → Make healthcare assistance accessible to non-English speakers.
- **Voice-Based Interaction** → Add speech-to-text and text-to-speech for hands-free health consultations.

- **Emergency Alerts** → Trigger notifications if symptoms suggest a critical condition.

4. Deployment Improvements

- **Persistent Hosting** → Migrate from temporary Colab/Streamlit deployments to cloud platforms (AWS, IBM Cloud, GCP, Azure).
- **Mobile Application** → Develop Android/iOS apps for accessibility on smartphones.
- **Offline Mode** → Lightweight version for low-connectivity regions.

5. Security & Compliance

- **End-to-End Encryption** → Ensure secure transmission of sensitive health data.
- **HIPAA/GDPR Compliance** → Align system with healthcare regulations for realworld deployment.
- **Role-Based Access Control** → Provide different access levels for patients, doctors, and administrators.

6. AI Model Enhancements

- **Smarter Prompt Engineering** → Predefined medical prompts for more consistent outputs.
- **Hybrid Model Usage** → Combine large models (accuracy) with lightweight models (speed).
- **Continuous Learning** → Enable the system to adapt based on user feedback and updated medical guidelines.